

**Unmanned Maritime Autonomy Architecture (UMAA)
Support Operations (SO)
Interface Control Document (ICD)
(UMAA-SPEC-SOICD)**

Version 6.0

6 June 2024

DISTRIBUTION STATEMENT A: Approved for public release. Distribution is unlimited.

Contents

1 Scope	5
1.1 Identification	5
1.2 Overview	5
1.3 Document Organization	7
2 Referenced Documents	8
3 Introduction to Data Model, Services, and Interfaces	9
3.1 Data Model	9
3.2 Definitions	9
3.3 Data Distribution Service (DDST TM)	9
3.4 Naming Conventions	10
3.5 Namespace Conventions	11
3.6 Cybersecurity	12
3.7 GUID algorithm	12
3.8 Large Collections	12
3.8.1 Necessary QoS	12
3.8.2 Creating Large Collections	12
3.8.3 Updating Large Collections	14
3.8.4 Removing an element from Large Collections	17
3.8.5 Specifying an Empty Large Collection	18
3.8.6 Large Set Types	18
3.8.7 Large List Types	19
3.9 Generalizations and Specializations	20
3.9.1 Creating a generalization/specialization	20
3.9.2 Updating a generalization/specialization	21
3.9.3 Removing a generalization/specialization	22
4 Flow Control	24
4.1 Command / Response	24
4.1.1 High-Level Flow	26
4.1.2 Command Startup Sequence	27
4.1.2.1 Service Provider Startup Sequence	27
4.1.2.2 Service Consumer Startup Sequence	28
4.1.3 Command Execution Sequences	29
4.1.4 Command Start Sequence	29
4.1.4.1 Command Execution	30
4.1.4.2 Updating a Command	31
4.1.4.3 Command Execution Success	32
4.1.4.4 Command Execution Failure	33
4.1.4.5 Command Canceled	34
4.1.5 Command Cleanup	35
4.1.6 Command Shutdown Sequence	36
4.1.6.1 Service Provider Shutdown Sequence	36
4.1.6.2 Service Consumer Shutdown Sequence	37
4.2 Request / Reply	38
4.2.1 Request/Reply without Query Data	38
4.2.1.1 Service Provider Startup Sequence	39
4.2.1.2 Service Consumer Startup Sequence	40
4.2.1.3 Service Provider Shutdown	40
4.2.1.4 Service Consumer Shutdown	40
4.2.2 Request/Reply with Query Data	41
5 Support Operations (SO) Services and Interfaces	42
5.1 Services and Interfaces	42

5.1.1	HealthReport	42
5.1.1.1	reportHealth	42
5.1.2	LogReport	43
5.1.2.1	reportLogReport	43
5.2	Common Data Types	45
5.2.1	UCSMDEInterfaceSet	45
5.2.2	UMAACCommand	45
5.2.3	UMAAStatus	45
5.2.4	UMAACCommandStatusBase	46
5.2.5	UMAACCommandStatus	46
5.2.6	DateTime	46
5.2.7	IdentifierType	47
5.3	Enumerations	48
5.3.1	CommandStatusReasonEnumType	48
5.3.2	ErrorCodeEnumType	48
5.3.3	ErrorConditionEnumType	49
5.3.4	LogLevelEnumType	49
5.3.5	CommandStatusEnumType	49
5.4	Type Definitions	51
A	Appendices	52
A.1	Glossary	52
A.2	Acronyms	52

List of Figures

1	UMAA Functional Organization.	5
2	UMAA Services and Interfaces Example.	6
3	Services and Interfaces Exposed on the UMAA Data Bus.	9
4	Sequence Diagram for initialization of a Large Collection with 3 elements.	13
5	Sequence Diagram for initialization of a Large Collection with 3 elements.	14
6	Sequence Diagram for update of Large Collection.	15
7	Sequence Diagram for update of an element of a Large Collection multiple times.	16
8	Sequence Diagram for delete of element from Large Collection.	17
9	Sequence Diagram for initialization of an empty Large Collection.	18
10	Generalization/Specialization UML diagram.	20
11	Sequence diagram for creating a generalization/specialization.	21
12	Sequence diagram for updating a generalization/specialization.	22
13	Sequence diagram for removing a generalization/specialization.	23
14	State transitions of the <code>commandStatus</code> as commands are processed.	25
15	Valid <code>commandStatusReason</code> values for each <code>commandStatus</code> state transition. Entries marked with a (—) indicate that the state transition is invalid.	25
16	Sequence Diagram for the High-Level Description of a Command Execution.	26
17	Sequence Diagram for Command Startup.	27
18	Sequence Diagram for Command Startup for Service Providers.	28
19	Sequence Diagram for Command Startup for Service Consumers.	29
20	Sequence Diagram for the Start of a Command Execution.	30
21	Beginning Sequence Diagram for a Command Execution.	31
22	Sequence Diagram for Command Update.	32
23	Sequence Diagram for a Command That Completes Successfully.	33
24	Sequence Diagram for a Command That Fails due to Resource Failure.	33
25	Sequence Diagram for a Command That Times Out Before Completing.	34
26	Sequence Diagram for a Command That is Canceled by the Service Consumer Before the Service Provider can Complete It.	35
27	Sequence Diagram Showing Cleanup of the Bus When a Command Has Been Completed and the Service Consumer No Longer Wishes to Maintain the Commanded State.	36

28	Sequence Diagram for Command Shutdown	36
29	Sequence Diagram for Command Shutdown for Service Providers	37
30	Sequence Diagram for Command Shutdown for Service Consumers	38
31	Sequence Diagram for a Request/Reply for Report Data That Does Not Require any Specific Query Data	39
32	Sequence Diagram for Initialization of a Service Provider to Provide FunctionReportTypes	40
33	Sequence Diagram for Initialization of a Service Consumer to Request FunctionReportTypes	40
34	Sequence Diagram for Shutdown of a Service Provider	40
35	Sequence Diagram for Shutdown of a Service Consumer	41

List of Tables

1	Standards Documents	8
2	Government Documents	8
3	Service Requests and Associated Responses	10
4	LargeSetMetadata Structure Definition	18
5	Example FooReportTypeItemsSetElement Structure Definition	19
6	LargeListMetadata Structure Definition	19
7	Example FooReportTypeItemsListElement Structure Definition	19
8	HealthReport Operations	42
9	HealthReportType Message Definition	43
10	LogReport Operations	43
11	LogReportType Message Definition	43
12	UCSMDEInterfaceSet Structure Definition	45
13	UMAACCommand Structure Definition	45
14	UMAAStatus Structure Definition	45
15	UMAACCommandStatusBase Structure Definition	46
16	UMAACCommandStatus Structure Definition	46
17	DateTime Structure Definition	46
18	IdentifierType Structure Definition	47
19	CommandStatusReasonEnumType Enumeration	48
20	ErrorCodeEnumType Enumeration	48
21	ErrorConditionEnumType Enumeration	49
22	LogLevelEnumType Enumeration	49
23	CommandStatusEnumType Enumeration	50
24	Type Definitions	51

1 Scope

1.1 Identification

This document defines a set of services and interfaces as part of the Unmanned Maritime Autonomy Architecture (UMAA). The services and corresponding interfaces covered in this ICD encompass the functionality to provide support operations for an Unmanned Maritime Vehicle (UMV) (surface or undersea). As such, it provides services used across functional boundaries of UMAA ICD's such as logging, supporting startup and shutdown, providing emissions services, and resource control. This document is generated automatically from data models that define its services and interfaces as part of the Unmanned Systems (UxS) Control Segment (UCS) Architecture as extended by UMAA to provide autonomy services for unmanned vehicles.

To put each ICD in context of the UMAA Architecture Design Description (ADD), the UMAA functional decomposition mapping to UMAA ICDs is shown in Figure 1.




Figure 1: UMAA Functional Organization.

1.2 Overview

The fundamental purpose of UMAA is to promote the development of common, modular, and scalable software for unmanned vehicles that is independent of a particular autonomy implementation. Unmanned Maritime Systems (UMSs) consist of Command and Control (C2), one or more unmanned vehicles, and support equipment and software (e.g. recovery system, Post Mission Analysis applications). The scope of UMAA is focused on the autonomy that resides on-board the unmanned vehicle. This includes the autonomy for all classes of unmanned vehicles and must support varying levels of communication in mission (i.e., constant, intermittent, or none) with external systems. To enable modular development and upgrade of the functional capabilities of the on-board autonomy, UMAA defines eight high-level functions. These core functions include: Communications Operations, Engineering Operations, Maneuver Operations, Mission Management, Processing Operations, Sensor and Effector Operations, Situational Awareness, and Support Operations. In each of these areas, it is anticipated that new capabilities will be required to satisfy evolving Navy missions over time. UMAA seeks to define standard interfaces for these functions so that individual programs can leverage capabilities developed to these standard interfaces across programs that meet the standard interface specifications. Individual programs may group services and interfaces into components in

different ways to serve their particular vehicle's needs. However, the entire interface defined by UMAA will be required as defined in the ICDs for all services that are included in a component. This requirement is what enables autonomy software to be ported between heterogeneous UMAA-compliant vehicles with their disparate vendor-defined vehicle control interfaces without recoding to a vehicle-specific interface.

Support Operations provides capabilities for services that are shared across all of the other functional areas within UMAA. This support includes the ancillary infrastructure services and interfaces required to operate an unmanned vehicle. Standard interfaces are defined for startup and shutdown, logging of time-stamped event and attribute data, operational mode control (e.g. operational, simulation, maintenance, training), and resource control (i.e. managing which client is in control of a component).

Unlike the primary concerns of an unmanned vehicle system, such as propulsion control and sensor data processing, the support operations are not typically seen in an external view of the system. Standardization of these services provides a consistent way to manage internal modes and control hierarchies across platforms and programs.




Figure 2: UMAA Services and Interfaces Example.

1.3 Document Organization

This interface control document is organized as follows:

Section 1 – Scope: A brief purview of this document

Section 2 – Referenced Documents: A listing of associated government and non-government documents and standards

Section 3 – Introduction to Data Model, Services, and Interfaces: A description of the common data model across all services and interfaces

Section 4 – Flow Control: A description of different flow control patterns used throughout UMAA

Section 5 – Support Operations (SO) Services and Interfaces: A description of specific services and interfaces for this ICD

2 Referenced Documents

The documents in the following table were used in the creation of the UMAA interface design documents. Not all references may be applicable to this particular document.

Table 1: Standards Documents

Title	Release Date
A Universally Unique IDentifier (UUID) URN Namespace	July 2005
Data Distribution Service for Real-Time Systems Specification, Version 1.4	March 2015
Data Distribution Service Interoperability Wire Protocol (DDSI-RTPS), Version 2.3	April 2019
Object Management Group Interface Definition Language Specification (IDL)	March 2018
Extensible and Dynamic Topic Types for DDS, Version 1.3	February 2020
UAS Control Segment (UCS) Architecture, Architecture Description, Version 2.4	27 March 2015
UCS Architecture, Conformance Specification, Version 2.2	27 September 2014
UCS-SPEC-MODEL v3.4 Enterprise Architect Model	27 March 2015
UCS Architecture, Architecture Technical Governance, Version 2.5	27 March 2015
System Modeling Language Specification, Version 1.5	May 2017
Unified Modeling Language Specification, Version 2.5.1	December 2017
Interface Definition Language (IDL), Version 4.2	March 2018
U.S. Department Of Homeland Security, United States Coast Guard "Navigation Rules International-Inland" COMDTINST M16672.2D	March 1999
IEEE 1003.1-2017 - IEEE Standard for Information Technology—Portable Operating System Interface (POSIX(R)) Base Specifications, Issue 7	December 2017
Guard, U. C. (2018). Navigation Rules and Regulations Handbook: International—Inland. Simon and Schuster.	June 2018
Department of Defense Interface Standard: Joint Military Symbology (MIL-STD-2525D Appendix A)	10 June 2014
DOD Dictionary of Military and Associated Terms	August 2018

Table 2: Government Documents

Title	Release Date
Unmanned Maritime Autonomy Architecture (UMAA) Architecture Design Description (ADD), Version 1.0	January 2019
Manual for the Submission of Oceanographic Data Collected by Unmanned Undersea Vehicles (UUVs)	October 2018

3 Introduction to Data Model, Services, and Interfaces

3.1 Data Model

A common data model is at the heart of UMAA. The common data model describes the entities that represent system state data, the attributes of those entities and relationships between those entities. This is a "data at rest" view of system-level information. It also contains data classes that define types of messages that will be produced by components, or a "data in motion" view of system-level information.

The common data model and coordinated service interfaces are described in a Unified Modeling Language (UMLTM) modeling tool and are represented as UMLTM class diagrams. Interface definition source code for messages/topics and other interface definition products and documentation will be automatically generated from the common data model so that they are consistent with the data model and to ensure that delivered software matches its interface specification.

The data model is maintained as a Multi-Domain Extension (MDE) to the UCS Architecture and will be maintained under configuration control by the UMAA Board as UCSMDE and will be incrementally integrated into the core UCS standard. Section 5 content is automatically generated from this data model, as are other automated products such as IDL that are used for automated code generation.

3.2 Definitions

UMAA ICDs follow the UCS terminology definitions found in the UCS Architecture Description v2.4. The normative (required) implementation to satisfy the requirements of a UMAA ICD is to provide service and interface specification compliance. Components may group services and required interfaces in any manner so long as every service meets its interface specifications. Figure 3 shows a particular grouping of services into components. The interfaces are represented by the blue and green lines and may equate to one or more independent input and output interfaces for each service. The implementation of the service into software components is left up to the individual system development. Given this context, section 5 correspondingly defines services with their interfaces and not components.

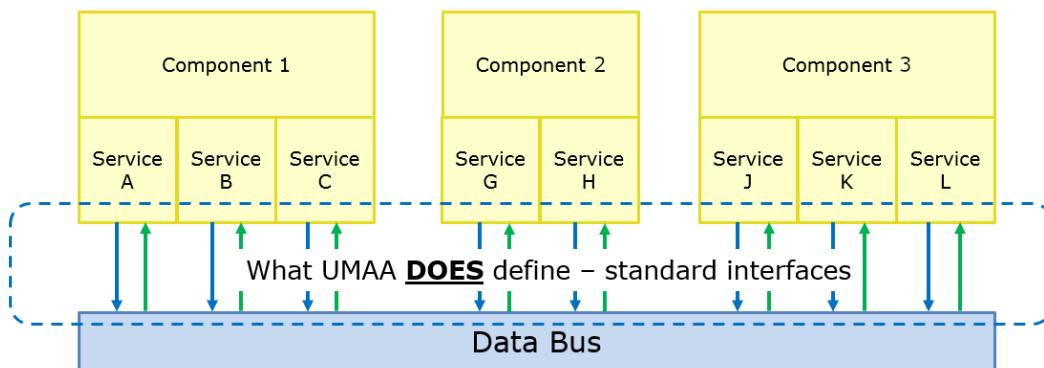


Figure 3: Services and Interfaces Exposed on the UMAA Data Bus.

Services may use other services within this ICD, or in other UMAA defined ICDs, to provide their capability. Additionally, components for acquisition and development may span multiple ICDs. An example of this would be a commercial radar that provides both status and control of the unit via the radar's software Application Programming Interface (API).

3.3 Data Distribution Service (DDSTM)

The data bus supporting autonomy messaging (as seen in Figure 3) is implemented via DDSTM. DDS is a middleware protocol and API standard for data-centric connectivity from the Object Management Group (OMG). It integrates the components of a system together, providing low-latency data connectivity, extreme reliability, and a scalable architecture. In a distributed system, middleware is the software layer that lies between the operating system and applications. It enables the various system components to more easily communicate and share data. It simplifies the development of distributed systems by letting software developers focus on the specific purpose of their applications rather than the mechanics of passing information between applications and systems. The DDS specification is fully described in free reference material on the OMG website and there are both open source and commercially available implementations.

3.4 Naming Conventions

UMAA services are modeled within the UCS Architecture under the Multi-Domain Extension (MDE). The UCS Architecture uses SoaML concepts of participant, serviceInterface, service port, and request port to describe the interfaces that make up a service and show how the service is used. Each service defines the capability it provides as well as required interfaces. Each interface consists of an operation that accepts a single message (A SoaML MessageType). In SoaML, a MessageType is defined as a unit of information exchanged between participant Request and Service ports via ServiceInterfaces. Instances of a MessageType are passed as parameters in ServiceInterface operations. (Reference: [UCS Architecture, Architecture Technical Governance](#))

To promote commonality across service definitions, a common way of naming services and their sets of operations and messages has been adopted for defining services within UCS-MDE. The convention uses the Service Base Name <SBN> and an optional Function Name [FN] to derive all service names and their associated operations and messages. As this is meant to be a guide, services might not include all of the defined operations and messages and their names might not follow the convention where a more appropriate name adds clarity.

Furthermore, services in UMAA are not required to be defined as indicated in Table 3 when all parts of the service capabilities are required for the service to be meaningful (such as ResourceAllocation).

Additionally, note that for UMAA not all operations defined in UCS-MDE result in a message being published to the DDS bus, e.g., since DDS uses publish/subscribe, most query operations result in a subscription to a topic and do not actually publish the associated request message. In the case of cancel commands, there is no associated implementation of the cancel<SBN>[FN]CommandStatus as it is just the intrinsic response of the DDS dispose function; so, it is essentially a NOOP (no operation) in implementation. The conventions used to define UCS-MDE services are as follows:

Service Name

- <SBN>[FN]Config
- <SBN>[FN]Control
- <SBN>[FN]Specs
- <SBN>[FN]Status OR Report

where the SBN should be descriptive of the task or information provided by the service. Note that the FN is optional and only included if needed to clarify the function of the service. The suffixes Status and Report are interchangeable. If a "Report" is a more appropriate description of the service, it can be used in lieu of "Status".

Table 3: Service Requests and Associated Responses

	Service Requests (Inputs)	Service Responses (Outputs)
Config	set<SBN>[FN]Config	report<SBN>[FN]ConfigCommandStatus
	query<SBN>[FN]ConfigAck	report<SBN>[FN]ConfigAck
	query<SBN>[FN]Config	report<SBN>[FN]Config
	cancel<SBN>[FN]Config	report<SBN>[FN]CancelConfigCommandStatus
	query<SBN>[FN]ConfigExecutionStatus	report<SBN>[FN]ConfigExecutionStatus
Control	set<SBN>[FN]	report<SBN>[FN]CommandStatus
	query<SBN>[FN]CommandAck	report<SBN>[FN]CommandAck
	cancel<SBN>[FN]Command	report<SBN>[FN]CancelCommandStatus
	query<SBN>[FN]ExecutionStatus	report<SBN>[FN]ExecutionStatus
Specs	query<SBN>[FN]Specs	report<SBN>[FN]Specs
Status OR Report	query<SBN>[FN]	report<SBN>[FN]

```

set<SBN>[FN]Config:<SBN>[FN]ConfigCommandType
query<SBN>[FN]Config:<SBN>[FN]ConfigRequestType1
set<SBN>[FN]:<SBN>[FN]CommandType
query<SBN>[FN]CommandAck:<SBN>[FN]CommandAckRequestType1
cancel<SBN>[FN]Command:<SBN>[FN]CancelCommandType1
cancel<SBN>[FN]Config:<SBN>[FN]CancelConfigType1
query<SBN>[FN]ExecutionStatus:<SBN>[FN]ExecutionStatusRequestType1
query<SBN>[FN]ConfigExecutionStatus:<SBN>[FN]ConfigExecutionStatusRequestType1
query<SBN>[FN]ConfigAck:<SBN>[FN]ConfigAckRequestType1
query<SBN>[FN]Specs:<SBN>[FN]SpecsRequestType1
query<SBN>[FN]:<SBN>[FN]RequestType1 2

```

Service Responses (operation:message)

```

report<SBN>[FN]ConfigCommandStatus:<SBN>[FN]ConfigCommandStatusType
report<SBN>[FN]Config:<SBN>[FN]ConfigReportType
report<SBN>[FN]ConfigAck:<SBN>[FN]ConfigAckReportType
report<SBN>[FN]CommandStatus:<SBN>[FN]CommandStatusType
report<SBN>[FN]CommandAck:<SBN>[FN]CommandAckReportType
report<SBN>[FN]CancelCommandStatus:<SBN>[FN]CancelCommandStatusType1
report<SBN>[FN]CancelConfigCommandStatus:<SBN>[FN]CancelConfigCommandStatusType1
report<SBN>[FN]ExecutionStatus:<SBN>[FN]ExecutionStatusReportType
report<SBN>[FN]ConfigExecutionStatus:<SBN>[FN]ConfigExecutionStatusReportType
report<SBN>[FN]Specs:<SBN>[FN]SpecsReportType
report<SBN>[FN]:<SBN>[FN]ReportType

```

where,

- Config (Configuration) Command/Report – This is the setup of a resource for operation of a particular task. Attributes may be static or variable. Examples include: maximum RPM allowed, operational sonar frequency range allowed, and maximum allowable radio transmit power.
- Command Status – This is the current state of a particular command (either control or configuration).
- Command – This is the ability to influence or direct the behavior of a resource during operation of a particular task. Attributes are variable. Examples include a vehicle's speed, engine RPM, antenna raising/lowering, and controlling a light or gong.
- Command Ack (Acknowledgement) Report – This is the command currently being executed.
- Cancel – This is the ability to cancel a particular command that has been issued.
- Execution Status Report – This is the status related to executing a particular command. Examples associated with a waypoint command include cross track error, time to achieve, and distance remaining.
- Specs (Specifications) Report – Provides a detailed description of a resource and/or its capabilities and constraints. Attributes are static. Examples include: maximum RPM of a motor, minimum frequency of a passive sonar sensor, length of the unmanned vehicle, and cycle time of a radar.
- Report – This is the current information being provided by a resource. Examples include vehicle speed, rudder angle, current waypoint, and contact bearing.

3.5 Namespace Conventions

Each UMAA service and the messages under the service can be accessed through their appropriate UMAA namespace. The namespace reflects the mapping of a specific service to its parent ICD, and the parent ICD's mapping to the overall UMAA Design Description. For example:

Access the Primitive Driver Control service under Maneuver Operations:

¹These message types are required for UCS model rules of construction, but are not implemented as messages in the UMAA specification.

²At this time, there are no Requests in the specification. This will be the message format when Requests have been added.

UMAA::MO::PrimitiveDriverControl
Access the ContactReport Service under Situational Awareness:
UMAA::SA::ContactReport

The UMAA model uses common data types that are re-used through the model to define service interface topics, interface topics, and other common data topics. These data types are not intended to be directly utilized but, for reference, they can be accessed in the same manner:

Access the common UMAA Status Message Fields:
UMAA::UMAAStatus
Access the common UMAA GeoPosition2D (i.e., latitude and longitude) structure:
UMAA::Common::Measurement::GeoPosition2D

3.6 Cybersecurity

The UMAA standard addressed in this ICD is independent from defining specific measures to achieve Cybersecurity compliance. This UMAA ICD does not preclude the incorporation of security measures, nor does it imply or guarantee any level of Cybersecurity within a system. Cybersecurity compliance will be performed on a program-specific basis and compliance testing is outside the scope of UMAA.

3.7 GUID algorithm

The UMAA standard utilizes the Globally Unique IDentifier (GUID), conforming to the variant defined in RFC 4122 (variant value of 2). Generators of GUIDs may generate GUIDs of any valid, RFC 4122-defined version that is appropriate for their specific use case and requirements. (Reference: [A Universally Unique IDentifier \(UUID\) URN Namespace](#))

3.8 Large Collections

The UMAA standard defines Large Collections, which are collections of decoupled but related data. Large Collections provide the ability to update one or more elements of the collection without republishing the entire collection to the DDS bus. This avoids two problems related to using an unbounded sequence type in a DDS message: 1) resource consumption growing as the collection is appended to or updated, and 2) DDS implementation-specific limitations on unbounded sequences. There are two implementations of a Large Collection: the Large Set (unordered) and the Large List (ordered).

In both Large Collection implementations, there are two important abstractions: the collection metadata and collection element type. Because Large Collections are specific to the UMAA PSM, the type definitions for the collection metadata and collection element are not part of MDE, and the IDL definitions of these types are generated separately. A particular UMAA message that has a Large Collection attribute will reference the metadata type (LargeSetMetadata or LargeListMetadata). The collection element type is defined under the same namespace as the message that uses it, and follows the naming pattern <parent message name><attribute name><collection type>Element. Each element of the collection is published as a separate message on the DDS bus, and can be tracked back to their related collection using the setID or listID. Users can also trace an element in a set to the source attribute (a NumericGUID) of the Service Provider that generated the report with this set using the collection metadata.

3.8.1 Necessary QoS

To achieve the Large Collection consistency in the update process described below, ordering of samples on the collection element type topic is necessary. Therefore, publishers and subscribers to the collection element type topic must use the PRESENTATION QoS policy with an access_scope of DDS_TOPIC_PRESENTATION_QOS and ordered_access.

Note that Large Collection Metadata and Elements are sent on separate DDS topics. DDS QoS does not guarantee ordering across topics. For this reason, implementations must be able to handle cases where elements arrive before or after the associated metadata. Memory must be allocated to await the proper metadata and associated elements.

3.8.2 Creating Large Collections

To create a large collection, a series of element messages and a metadata message must be sent from one DDS participant (the sender) to another (the receiver). The messages should be buffered on the receiving side until a synchronization point is

reached which indicates an atomic update. That is, when both a metadata message and an element message corresponding by list ID, timestamp, and last element ID have been received, yield a complete collection. Figure 4 shows the sequence of exchanges to establish a collection with 3 elements.

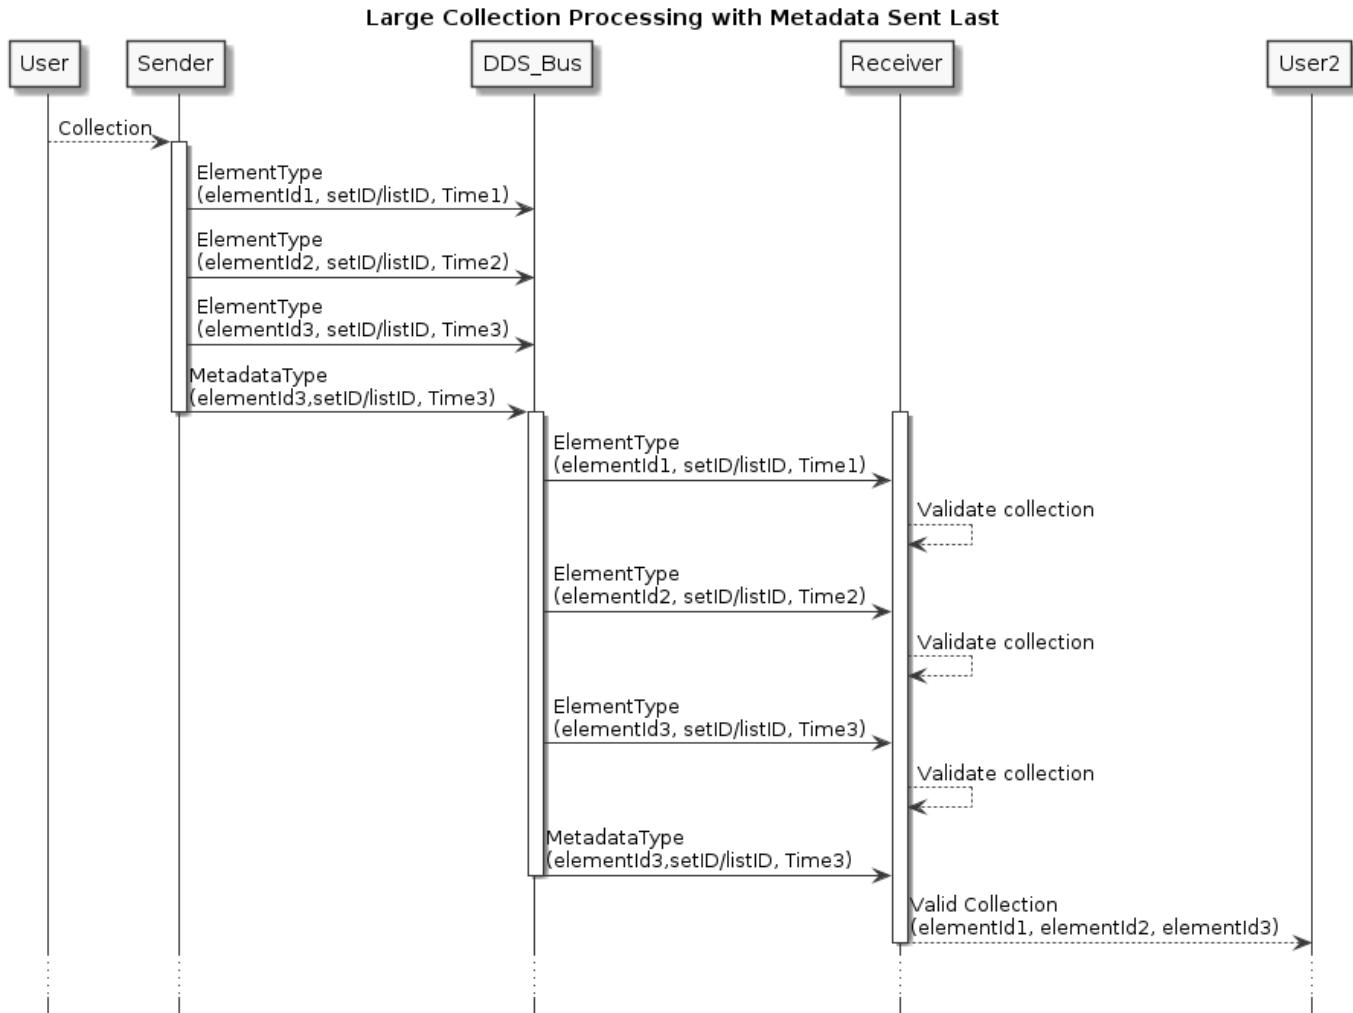


Figure 4: Sequence Diagram for initialization of a Large Collection with 3 elements.

The same collection could be established where the element data arrives after the metadata, creating the same list as depicted in figure 5.




Figure 5: Sequence Diagram for initialization of a Large Collection with 3 elements.

3.8.3 Updating Large Collections

When elements of the collection are updated, the metadata must be updated as well to signal a change in the set. The `updateElementID` is updated to match the `elementID` of the element whose reception signals the end of the atomic update of the collection. Because of the requirement of an ordered topic described above, this will be the element that is updated last chronologically. The metadata `updateElementTimestamp` must be updated to the timestamp of the same element that signals the end of the update.

The set can be updated as a batch (multiple elements in a single "update cycle," as determined by the provider). This allows for a coarse synchronization: data elements that do not match the metadata `updateElementID` and `updateElementTimestamp` can be assumed to be part of an in-progress update cycle. Consumers can choose to immediately act on those data individually or wait until the matching element is received to signal that the complete update cycle has finished and consider the set as a whole. Note that the coarseness of synchronization is service-dependent: in some cases an intermediate view of a collection update may be logically incorrect to act upon.

Figure 6 shows the sequence of exchanges to update a collection of 3 elements and add a 4th element.

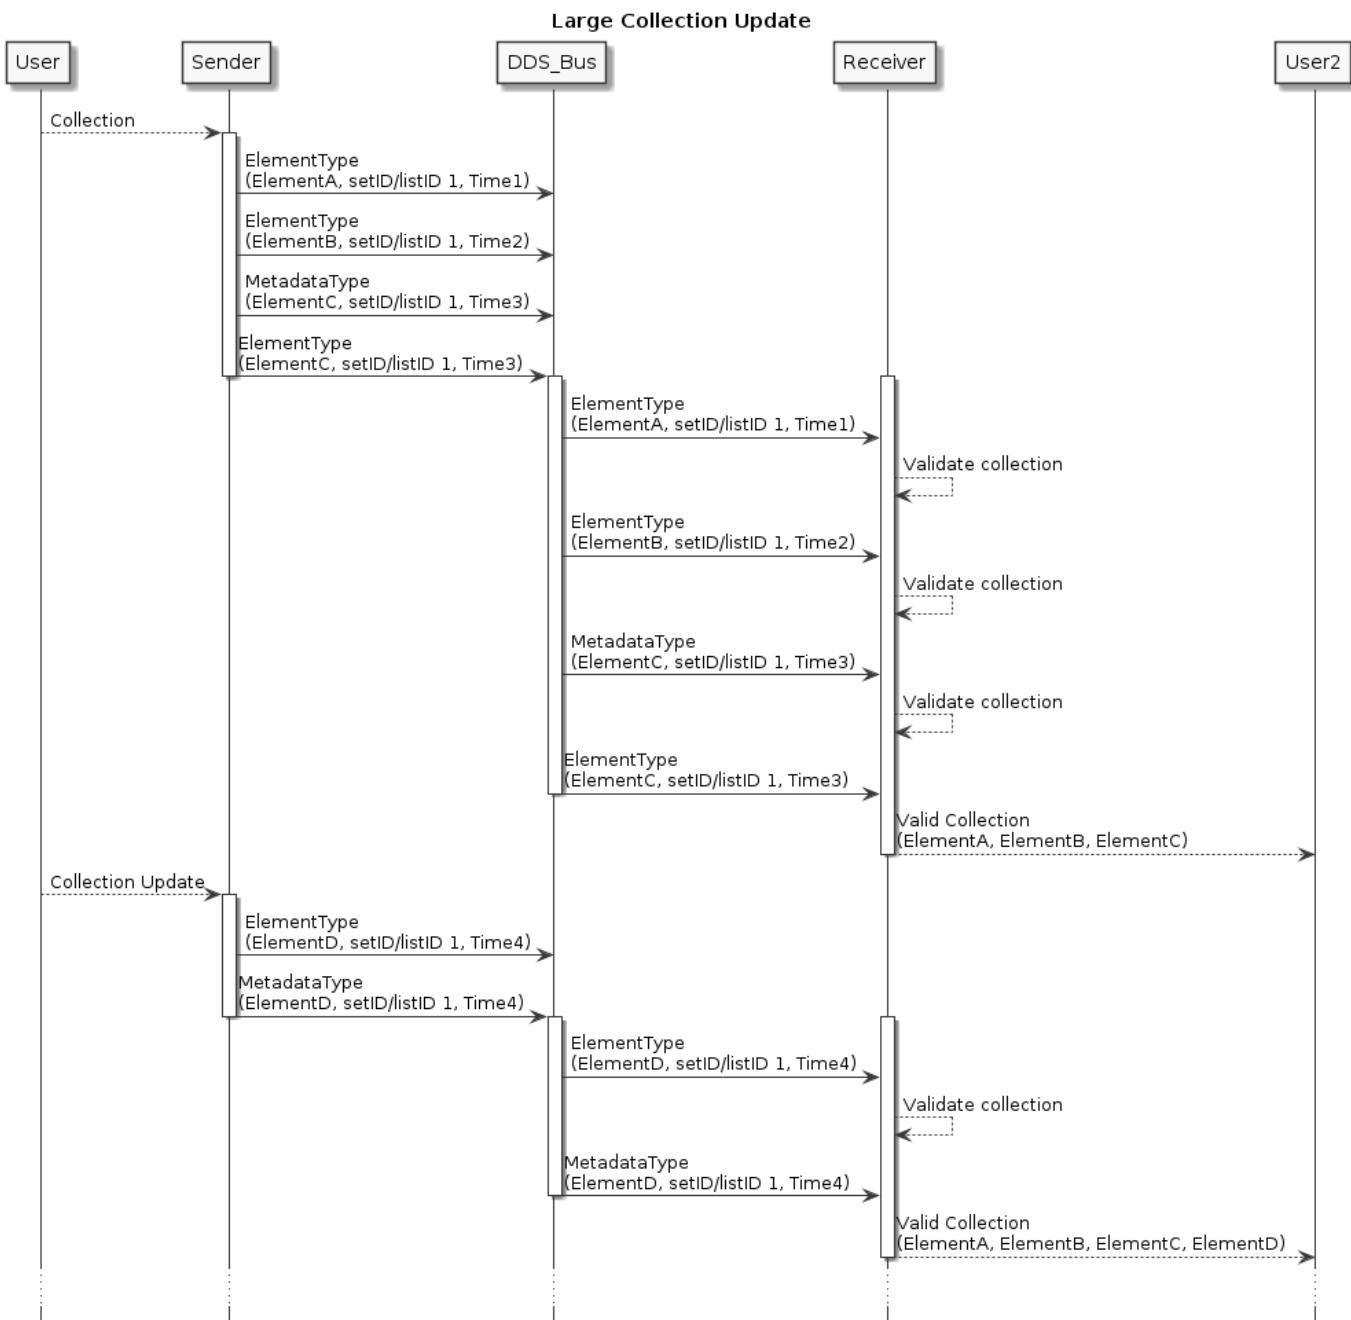


Figure 6: Sequence Diagram for update of Large Collection.

Figure 7 shows the sequence of exchanges to update an element of a collection multiple times.




Figure 7: Sequence Diagram for update of an element of a Large Collection multiple times.

3.8.4 Removing an element from Large Collections

To remove an element from a collection, dispose of the element on the element topic and re-publish the metadata. Multiple deletes and inserts can happen for a single metadata update. In the case where the final element of the collection is deleted, the updateElementTimestamp should be unset in the metadata.

Figure 8 shows the sequence of exchanges to delete an element from a Large Collection.




Figure 8: Sequence Diagram for delete of element from Large Collection.

For Large Lists, it may be necessary to update the nextElementID references during delete operations to ensure that the list is still valid. This would cause multiple element messages to be sent along with updated metadata.

3.8.5 Specifying an Empty Large Collection

A particular Large Collection can be empty during initial creation. This is indicated by publishing metadata with a `size` of zero and an `updateElementID` set to the Nil UUID. As specified in section 4.1.7 of the referenced document "A Universally Unique Identifier (UUID) URN Namespace", this is a "special form of UUID that is specified to have all 128 bits set to zero".

Figure 9 shows the sequence of exchanges to establish an initially empty Large Collection.




Figure 9: Sequence Diagram for initialization of an empty Large Collection.

3.8.6 Large Set Types

The following details the LargeSetMetadata structure:

Table 4: LargeSetMetadata Structure Definition

Attribute Name	Attribute Type	Attribute Description
setID	NumericGUID	Identifies the Large Set instance this metadata relates to.
updateElementID	NumericGUID	This field references the element ID of the set element whose reception signals the end of an atomic update to this set. This elementID must be used in conjunction with the updateElementTimestamp below to fully identify when the atomic update has completed and the set is stable.
updateElementTimestamp†	DateTime	This field identifies the elementTimestamp of the element, referenced above by updateElementID, that signals the end of an atomic update to this set. This field will be empty in the event that the element update results from a DDS dispose.
size	LargeCollectionSize	Indicates the number of elements associated with this set after the atomic update is complete.

An example element type is shown below, where a `FooReportType` message has a Large Set attribute called "items" whose type is `BarType`

Table 5: Example FooReportTypeItemsSetElement Structure Definition

Attribute Name	Attribute Type	Attribute Description
element	BarType	The value of the set element.
setID*	NumericGUID	Identifies the Large Set instance this element relates to.
elementID*	NumericGUID	Uniquely identifies this element within the set and across all large collection elements that currently exist on the DDS bus.
elementTimestamp	DateTime	The timestamp of this element.

3.8.7 Large List Types

The following details the LargeListMetadata structure:

Table 6: LargeListMetadata Structure Definition

Attribute Name	Attribute Type	Attribute Description
listID	NumericGUID	Identifies the Large List instance this metadata relates to.
updateElementID	NumericGUID	This field references the element ID of the list element whose reception signals the end of an atomic update to this list. This elementID must be used in conjunction with the updateElementTimestamp below to fully identify when the atomic update has completed and the list is stable.
updateElementTimestamp†	DateTime	This field identifies the elementTimestamp of the element, referenced above by updateElementID, that signals the end of an atomic update to this list. This field will be empty in the event that the element update results from a DDS dispose.
startingElementID	NumericGUID	This field identifies the list element, tying to its elementID, that is sequentially first in the list. This is provided for convenience when iterating through the linked list using the nextElementID field.
size	LargeCollectionSize	Indicates the number of elements associated with this set after the atomic update is complete.

An example element type is shown below, where a `FooReportType` message has a Large List attribute called "items" whose type is `BarType`

Table 7: Example FooReportTypeItemsListElement Structure Definition

Attribute Name	Attribute Type	Attribute Description
element	BarType	The value of the list element.
listID*	NumericGUID	Identifies the Large List instance this element relates to.
elementID*	NumericGUID	Uniquely identifies this element within the list and across all large collection elements that currently exist on the DDS bus.
elementTimestamp	DateTime	The timestamp of this element.

Attribute Name	Attribute Type	Attribute Description
nextElementID†	NumericGUID	This field references to the elementID of the element that logically follows this element in the linked list. This is empty if this element is sequentially last.

3.9 Generalizations and Specializations

The UMAA standard makes use of generalization/specialization relationships when defining data types. The generalization/specialization relationship is one where a generalization data structure is defined to contain attributes that are common across some entity and specialization data structures are defined to contain attributes that are specific to a particular type of that entity. This relationship can be modeled as inheritance in UML as shown below.

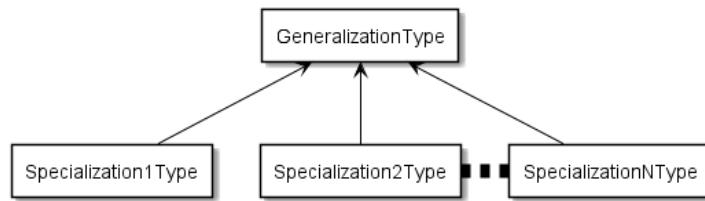


Figure 10: Generalization/Specialization UML diagram.

When the data type of an attribute within a message is a generalization, it is defined to be that generalization plus the data type of one of its specializations. In order to support this relationship, the generalization data structure and its specialization data structure are published to separate topics along with additional metadata linking the two topics. Specifically, the generalization data structure includes: specializationTopic, specializationID, and specializationTimestamp; and the specialization data structure includes: specializationID and specializationTimestamp. The specializationTopic specifies the topic name of the particular specialization, and the specializationID and specializationTimestamp must be equivalent in each topic, respectively, in order to establish the generalization/specialization relationship.

3.9.1 Creating a generalization/specialization

To create a generalization/specialization, both the GeneralizationType and SpecializationType topics must be sent from one DDS participant (the sender) to another (the receiver). The topics should be buffered on the receiving side until a synchronization point is reached that indicates an atomic update.

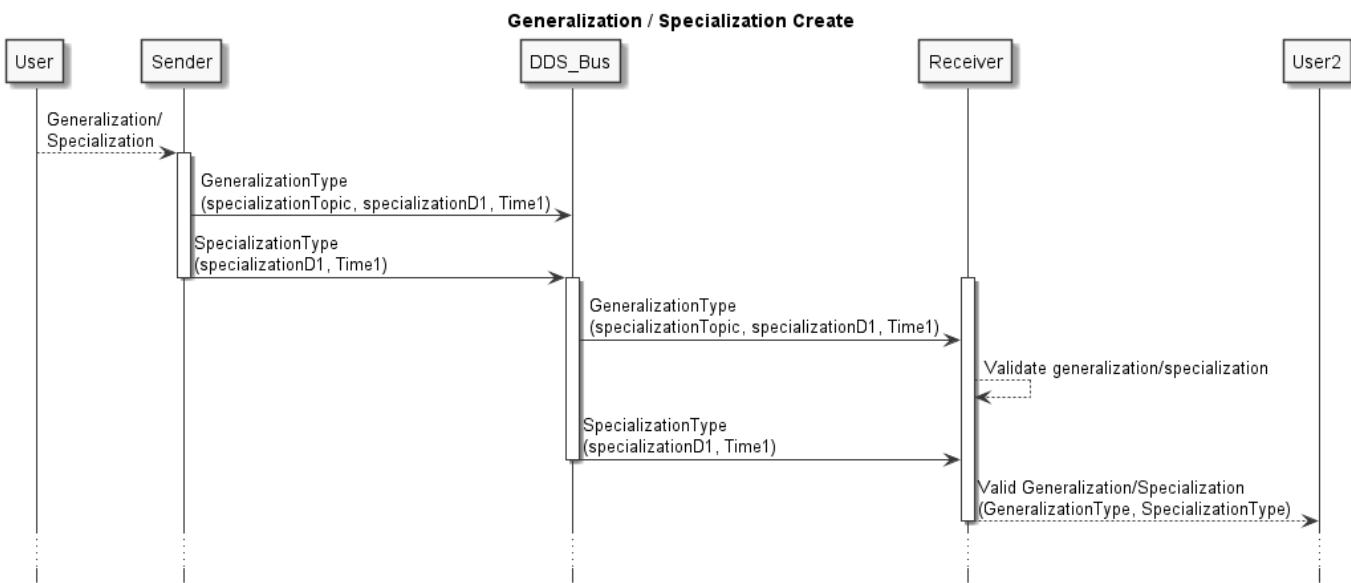


Figure 11: Sequence diagram for creating a generalization/specialization.

3.9.2 Updating a generalization/specialization

An update to a generalization/specialization can occur when there is a change in either data structure. In order for the update to be complete, the specializationTimestamp must be updated in both the GeneralizationType and the SpecializationType, and again they must be equal. Note that if a generalization/specialization exists within a large set or large list that their respective metadata must also be updated as defined in Section 3.8.




Figure 12: Sequence diagram for updating a generalization/specialization.

3.9.3 Removing a generalization/specialization

To remove a generalization/specialization, both topics must be disposed. Again, note that if a generalization/specialization exists within a large set or large list that their respective metadata must also be updated as defined in Section 3.8.




Figure 13: Sequence diagram for removing a generalization/specialization.

4 Flow Control

4.1 Command / Response

This section defines the flow of control for command/response over the DDS bus. A command/response controls a specific service. While the exact names and processes will depend on the specific service and command being executed, all command/responses in UMAA follow a similar pattern. A notional "Function" command **FunctionCommand** is used in the following examples. As will be described in subsequent paragraphs, DDS publish/subscribe methods are used in implementations to issue commands and responses.

To direct a **FunctionCommand** at a specific Service Provider, UMAA includes a **destination** GUID in all commands. A Service Provider is required to respond to all **FunctionCommands** where the **destination** is the same as the Service Provider's ID. The Service Consumer will also create a **sessionID** for the command when commanded. The **sessionID** is used to track the command execution as a key into other command-related messages. The **sessionID** must be unique across all **FunctionCommand** instances that are active (i.e. currently on the DDS bus), otherwise the Service Provider will consider the **FunctionCommand** to be a command update (see Section 4.1.4.2). Once a **FunctionCommand** is removed from the DDS bus as part of the Command Cleanup process (see Section 4.1.5), its **sessionID** may be reused for future commands without triggering a command update; therefore it is not necessary for a Service Provider to maintain a complete history of **sessionIDs**.

Service Provider and Service Consumer terminology in the following sections is adopted from the OMG Service-oriented architecture Modeling Language (SoaML).

To initialize, a Service Provider (controllable resource) subscribes to the **FunctionCommand** DDS topic. At startup or right before issuing a command, the Service Consumer (controlling resource) subscribes to the **FunctionCommandStatus** DDS topic. Optionally, the Service Consumer may also subscribe to the **FunctionCommandAckReport** to monitor which command is currently being executed, and the **FunctionExecutionStatusReport** (if defined for the Function service) that provides reporting on function-specific data status.

Both Service Providers and Service Consumers are required to recover or clean up any previous persisted commands on the bus during initialization.

To execute a command, the Service Consumer publishes a **FunctionCommandType** to the DDS bus. The Service Provider will be notified and will begin processing the request. During each phase of processing, the Service Provider will provide updates to the Service Consumer via published updates to a related **FunctionCommandStatus** topic. Command responses are correlated to their originating command via the **sessionID**. If a command with a duplicate **sessionID** is received, the Service Provider will regard this as a command update, and follow the flow control detailed in Section 4.1.4.2. Command status updates are provided in the command responses via the **commandStatus** field with additional details included in the **commandStatusReason** field. The Service Provider will also publish the current executing command to the **FunctionCommandAckReport** topic. When defined for the Function service, the Service Provider must also publish the **FunctionExecutionStatusReport** topic and update it as appropriate throughout the execution of the command.

The required state transitions for the **commandStatus** field are shown in Figure 14. Commands may complete normally, or they may terminate early due to failure (Section 4.1.4.4) or cancellation (Section 4.1.4.5). The state machine for a command can also be reset to **ISSUED** via a command update (Section 4.1.4.2). If there is not a self-transition indicated in the diagram, you cannot republish that state in a message. Every command must transition through the states as defined. For example, it is a violation to transition from **ISSUED** to **EXECUTING** without transitioning through **COMMANDED**. Even in the case where there is no logic executing between the **ISSUED** and **EXECUTING** states, the Service Provider is required to transition through **COMMANDED**. This ensures consistent behavior across different Service Providers, including those that do require the **COMMANDED** state.




Figure 14: State transitions of the `commandStatus` as commands are processed.

As described above, each time a command transitions to a new state, a `FunctionCommandStatus` message is published containing the updated `commandStatus` and a `commandStatusReason` that indicates why the state transition happened. The table below shows all valid `commandStatusReason` values for each `commandStatus` transition.

Starting State	Ending State					
	ISSUED	COMMANDED	EXECUTING	COMPLETED	FAILED	CANCELED
Initial State	SUCCEEDED	—	—	—	—	—
ISSUED	UPDATED	SUCCEEDED	—	—	VALIDATION_FAILED RESOURCE_FAILED INTERRUPTED TIMEOUT SERVICE_FAILED	CANCELED
COMMANDED	UPDATED	—	SUCCEEDED	—	RESOURCE_REJECTED INTERRUPTED TIMEOUT SERVICE_FAILED	CANCELED
EXECUTING	UPDATED	—	—	SUCCEEDED	OBJECTIVE_FAILED RESOURCE_FAILED INTERRUPTED TIMEOUT SERVICE_FAILED	CANCELED
COMPLETED	—	—	—	—	—	—
FAILED	—	—	—	—	—	—
CANCELED	—	—	—	—	—	—

Figure 15: Valid `commandStatusReason` values for each `commandStatus` state transition. Entries marked with a (—) indicate that the state transition is invalid.

In the following sections, the sequence diagrams demonstrate different exchanges between a Service Consumer and Service Provider. Within the diagrams, the dashed arrows represent implementation-specific communications that are outside of UMAA's scope. These sequence diagrams are just an example of one possible implementation. Other implementations may have different communication patterns between the Service Provider and the Resource or be implemented completely within the Service Provider process itself (no dependency on an external Resource). Likewise, the interactions between the User and Service Consumer may follow similar or different patterns. However, the UMAA-defined exchanges with the DDS bus between the Service Consumer and Service Provider must happen in the order shown within the sequence diagrams.

4.1.1 High-Level Flow

The high-level flow of a command sequence is shown in Figure 16 and can be described as follows:

1. The Command Startup Sequence is performed.
2. For each command to be executed:
 - (a) The Command Start Sequence is performed.
 - (b) The command is executed (sequence depends on the execution path, i.e., success, failure, or cancel).
 - (c) The Command Cleanup Sequence is performed.
3. The Command Shutdown Sequence is performed.

The **ref** blocks will be defined in later sequence diagrams. Note that the duration of the system execution for any particular **FunctionCommandType** is defined by the combination of the Service Provider(s) and Service Consumer(s) in the system and may not be identical to the overall system execution duration. For example, providers may only be available to execute certain commands during specific mission phases or when certain hardware is in specific configurations. This Command Startup Sequence is not required to happen during a system startup phase. The only requirement is that it must be completed by at least one Service Provider and one Service Consumer before any **FunctionCommandType** commands can be fully executed. Likewise, the Command Shutdown sequence may occur at any time the **FunctionCommandType** will no longer be supported. There is no requirement stating that the Command Shutdown Sequence only be performed during a system shutdown phase.




Figure 16: Sequence Diagram for the High-Level Description of a Command Execution.

4.1.2 Command Startup Sequence

As part of initialization both the Service Provider and Service Consumer are required to perform a startup sequence. This startup prepares the Service Provider to execute commands and the Service Consumer to request commands and monitor the progress of those requested commands.

The Service Provider and Service Consumer can initialize in any order. Commands will not be completely executed until both have completed their initialization. The sequence diagram is shown in Figure 17.

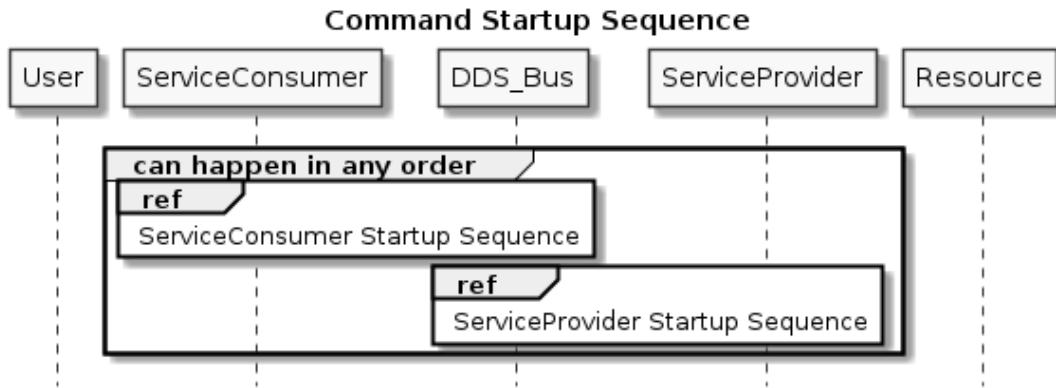


Figure 17: Sequence Diagram for Command Startup.

4.1.2.1 Service Provider Startup Sequence During startup, the Service Provider is required to register as a publisher to the `FunctionCommandStatus`, `FunctionCommandAckReport`, and (if defined for the Function service) the `FunctionExecutionStatusReport` topics.

The Service Provider is also required to subscribe to the `FunctionCommand` topic to be notified when new commands are published.

Finally, the Service Provider is required to handle any existing `FunctionCommandType` commands persisted on the DDS bus with the Service Provider's ID. For each command, if the Service Provider can and wishes to recover, it can continue to execute the command. To obtain the last published state of the command, the Service Provider must subscribe to the `FunctionCommandStatusType`. The Service Provider will continue following the normal status update sequence, picking up from the last status on the bus. If the Service Provider cannot or chooses not to continue processing the command, it must fail the command by publishing a `FunctionCommandStatus` with a `commandStatus` of `FAILED` and a `reason` of `SERVICE_FAILED`.

The Service Provider Startup sequence is shown in Figure 18.




Figure 18: Sequence Diagram for Command Startup for Service Providers.

4.1.2.2 Service Consumer Startup Sequence During startup, the Service Consumer is required to register as a publisher of the `FunctionCommandType`.

The Service Consumer is also required to subscribe to the `FunctionCommandStatusType` to monitor the execution of any published commands. The Service Consumer can optionally register for the `FunctionCommandAckReportType` and, if defined for the Function service, the `FunctionExecutionStatusReportType` if it desires to track additional status of the execution of commands.

Finally, the Service Consumer is required to handle any existing `FunctionCommandType` commands persisted on the DDS bus with this Service Consumer's ID. To find existing `FunctionCommandTypes` on the bus, it must first subscribe to the topic. If the Service Consumer can and wishes to recover, it can continue to monitor the execution of the command. If the Service Consumer cannot or chooses not to continue the execution of the command, it must cancel the command via the normal command cancel method.

The Service Consumer Startup sequence is shown in Figure 19.

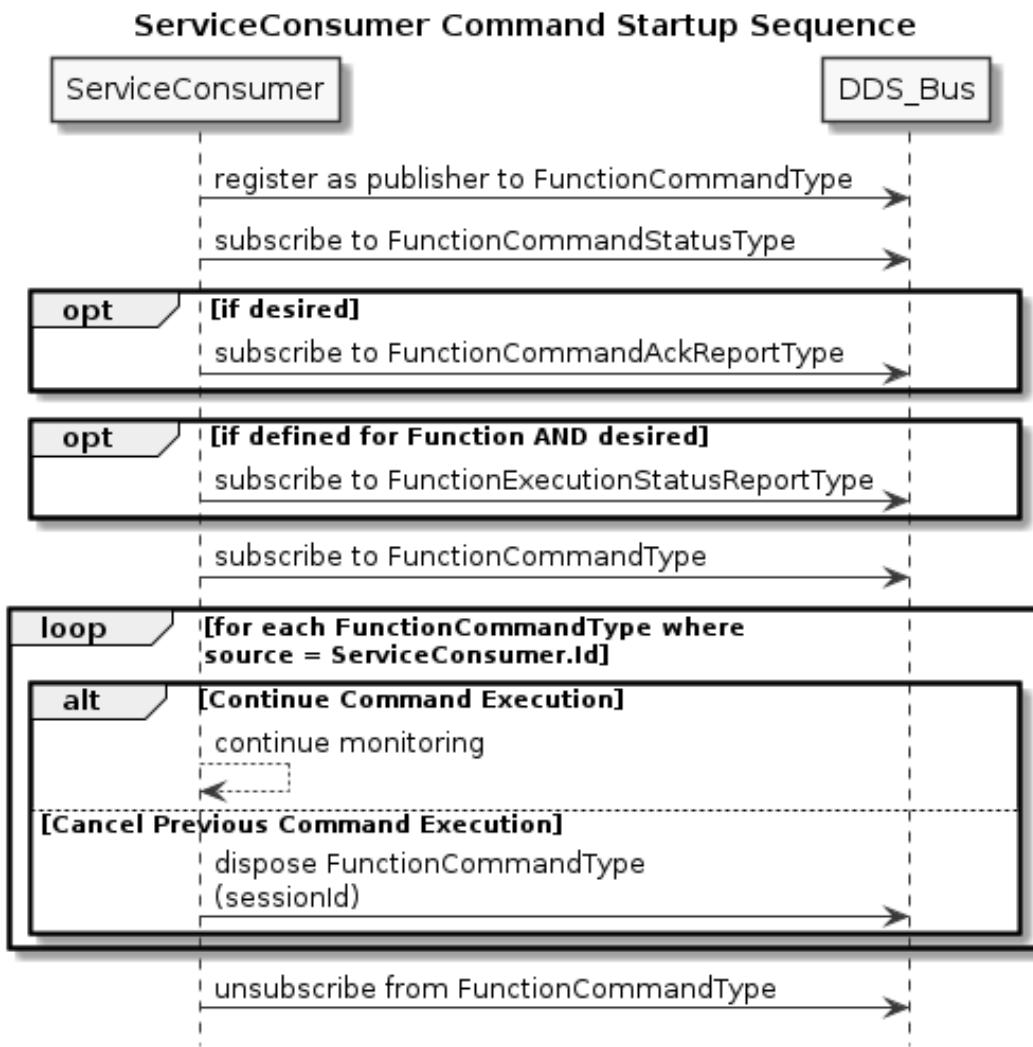


Figure 19: Sequence Diagram for Command Startup for Service Consumers.

4.1.3 Command Execution Sequences

Once both the Service Provider and Service Consumer have performed the startup sequence, the system is ready to begin issuing and executing commands.

4.1.4 Command Start Sequence

The initial start sequence to execute a single new command follows this pattern:

1. The User of the Service Consumer issues a request for a command to be executed.
2. The Service Consumer publishes the `FunctionCommandType` with a unique session ID, the source ID of the Service Consumer, and the destination ID of the desired Service Provider.
3. The Service Provider, upon notification of the new `FunctionCommandType`, publishes a new `FunctionCommandStatusType` with (1) the same session ID as the new `FunctionCommandType`, (2) the status of `ISSUED` and (3) the reason of `SUCCEEDED` to notify the Service Consumer it has received the new command.

The Command Start Sequence for a new command is shown in Figure 20. This pattern will be repeated each time a new command is requested. Note that the Command Start Sequence differs if the `FunctionCommandType` has a `sessionId` that matches another `FunctionCommandType` that currently exists on the DDS bus. This is considered a command update and detailed in Section 4.1.4.2.

After the Command Start Sequence, the sequence can take different paths depending on the actual execution of the command, detailed from Section 4.1.4.1 to Section 4.1.4.5, but they do not enumerate all of the possible execution paths. Other paths (e.g., an objective failing) will follow a similar pattern to other failures; all are required to follow the state diagram shown in Figure 14 and eventually end with the Command Cleanup Sequence (shown in Figure 27).




Figure 20: Sequence Diagram for the Start of a Command Execution.

4.1.4.1 Command Execution Once a Service Provider starts to process a command, the Command Execution sequence is:

1. The Service Provider publishes a `FunctionCommandAckReportType` with matching session ID and parameters as the `FunctionCommandType` it is starting to process.
2. The Service Provider performs any validation and negotiation with backing resources as necessary. Once the command is ready to be executed, the Service Provider publishes a `FunctionCommandStatusType` with a status `COMMANDED` and reason `SUCCEEDED` to notify the Service Consumer that the command has been validated and commanded to start execution.
3. Once the command has begun executing, the Service Provider publishes a `FunctionCommandStatusType` with a status `EXECUTING` and reason `SUCCEEDED` to notify the Service Consumer that the command has been validated and commanded to start.
4. If the Function has a defined `FunctionExecutionStatusReportType`, the Service Provider must publish a new instance with matching session ID as the associated `FunctionCommandType`. The `FunctionExecutionStatusReportType` must be updated by the Service Provider throughout the execution as dictated by the definitions of the command-specific attributes in the execution status report.

The command execution sequence is shown in Figure 21. This sequence holds until the command completes execution.

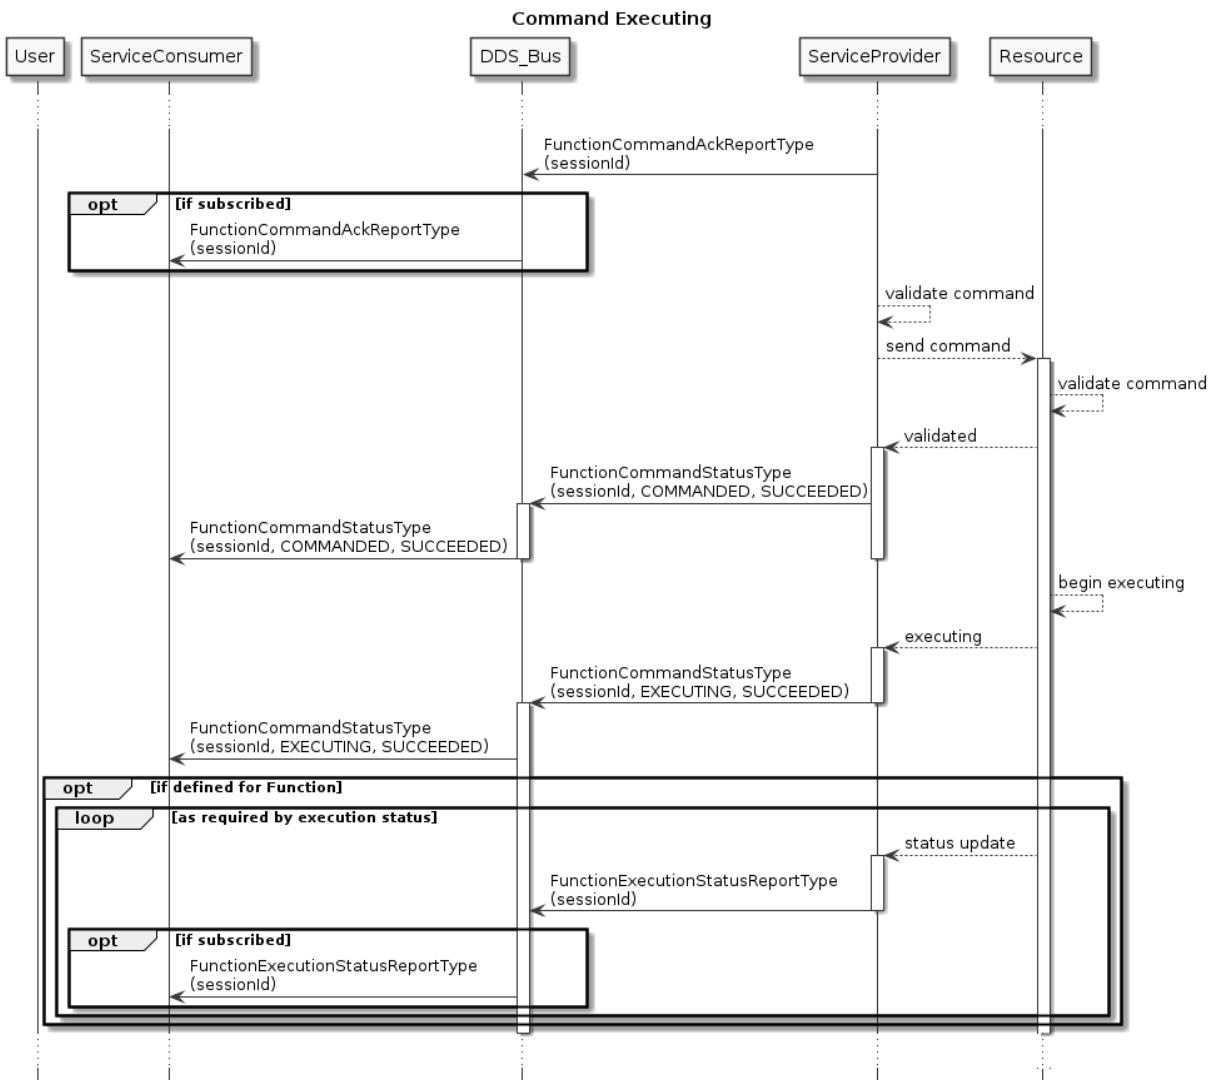


Figure 21: Beginning Sequence Diagram for a Command Execution.

The normal successful conclusion of a command being executed in some cases is initiated by the Service Consumer (an endless GlobalVector command concluded by canceling it) and in other cases is initiated by the Service Provider (a GlobalWaypoint commanded concluded by reaching the last waypoint). Unless otherwise explicitly stated, it is assumed the Service Provider will be able to identify the successful conclusion of a command. In the cases where commands are defined to be indeterminate the Service Consumer must cancel the command when the Service Consumer no longer desires the command to be executed.

4.1.4.2 Updating a Command An updated command is defined as a command with a source ID and session ID identical to the current command being processed by the Service Provider, but whose timestamp is newer than the current command. Only a command that is in a non-terminal state may be updated - otherwise, the Service Consumer must follow the normal command cleanup process and issue a new command with an updated unique session ID. If a command is in a terminal state, the Service Provider must ignore an update to that command.

When the Service Provider receives an updated command, it is required to take one of two possible actions:

1. If the current command is in a non-terminal state (ISSUED, COMMANDED, or EXECUTING), then the Service Provider publishes a FunctionCommandStatusType with a status ISSUED and reason UPDATED. The state machine then restarts and proceeds through the normal command flow detailed in 4.1.4. The Service Provider must consider the updated command as an entirely new command, resetting any internal state related to the command (e.g. a timer that keeps track of command timeout).

2. If the current command is in a terminal state (COMPLETED, CANCELED, or FAILED), then the updated command cannot be processed, and the Service Provider must publish a FunctionCommandStatusType with a status FAILED and follow the normal command cleanup process.

The flow control for command update is detailed below:

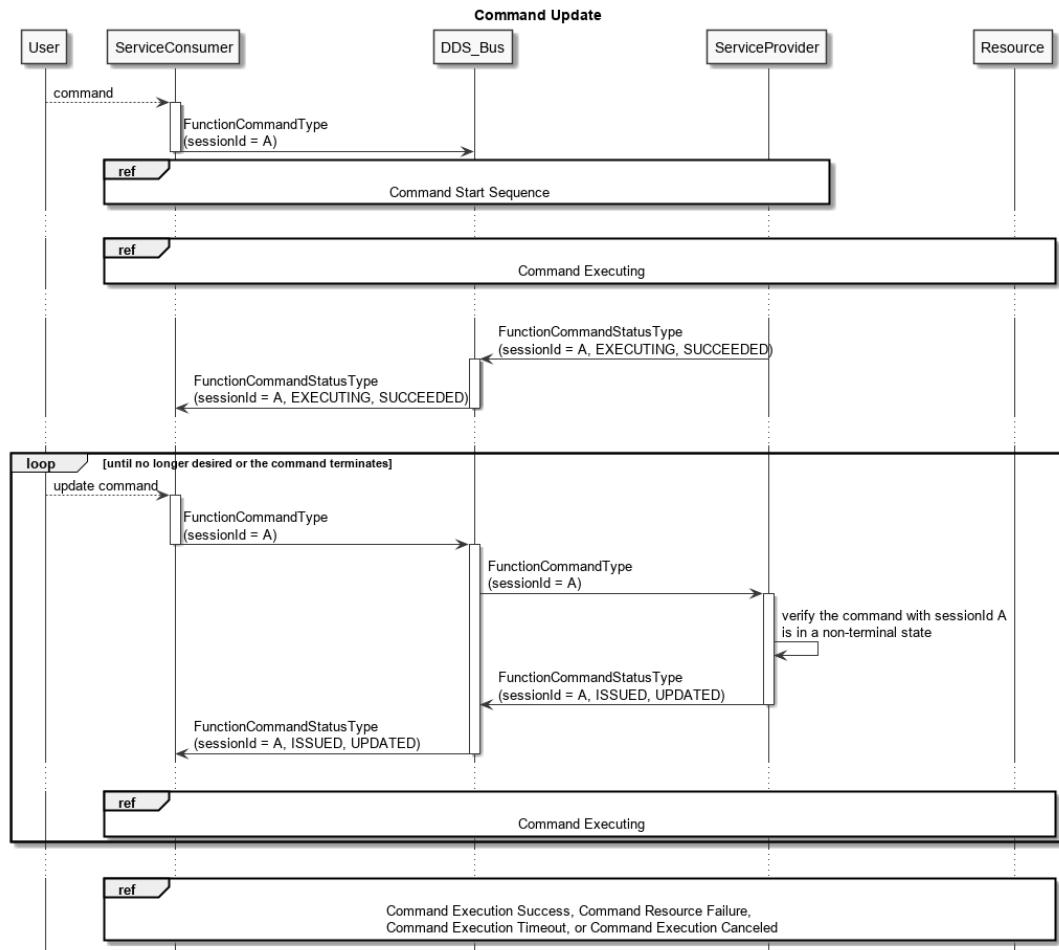


Figure 22: Sequence Diagram for Command Update.

4.1.4.3 Command Execution Success When the Service Provider determines a command has successfully completed, it must update the associated FunctionCommandStatusType with as status of COMPLETED and reason of SUCCEEDED. This signals to the Service Consumer that the command has completed successfully.

The Command Execution Success sequence is shown in Figure 23.

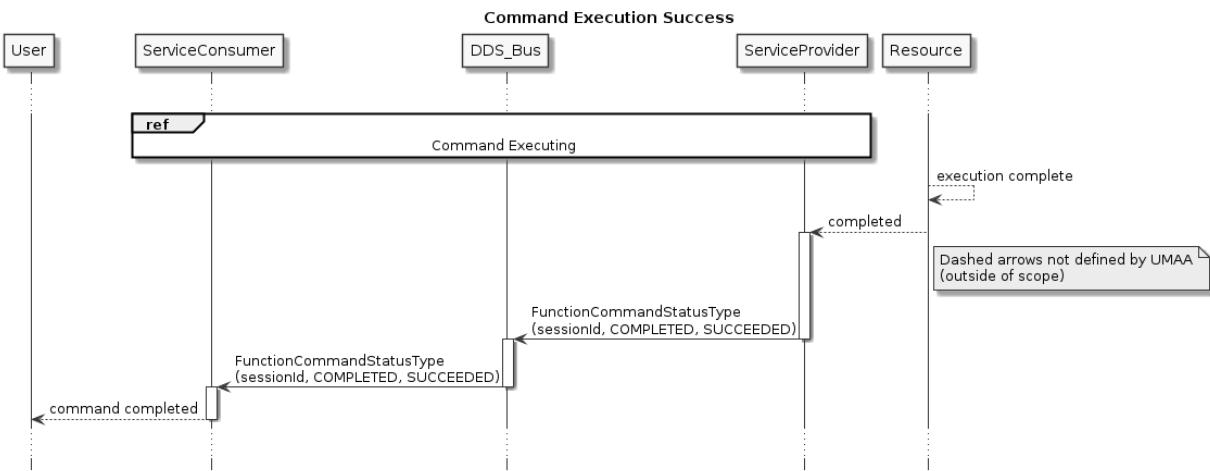


Figure 23: Sequence Diagram for a Command That Completes Successfully.

4.1.4.4 Command Execution Failure The command may fail to complete for any number of reasons including software errors, hardware failures, or unfavorable environmental conditions. The Service Provider may also reject a command for a number of reasons including inability to perform the task, malformed or out of range requests, or a command being interrupted by a higher priority process. In all cases, the Service Provider must publish a `FunctionCommandStatusType` with an identical `sessionId` as the originating `FunctionCommandType` with a status of `FAILED` and the reason that reflects the cause of the failure (`VALIDATION_FAILED`, `SERVICE_FAILED`, `OBJECTIVE_FAILED`, etc).

Figure 24 and Figure 25 provide examples where a command has failed.

In the first example, the backing Resource failed and the Service Provider is unable to communicate with it. In this case, the Service Provider will report a `FunctionCommandStatusType` with a status of `FAILED` and a reason of `RESOURCE_FAILED`. This is shown in Figure 24.

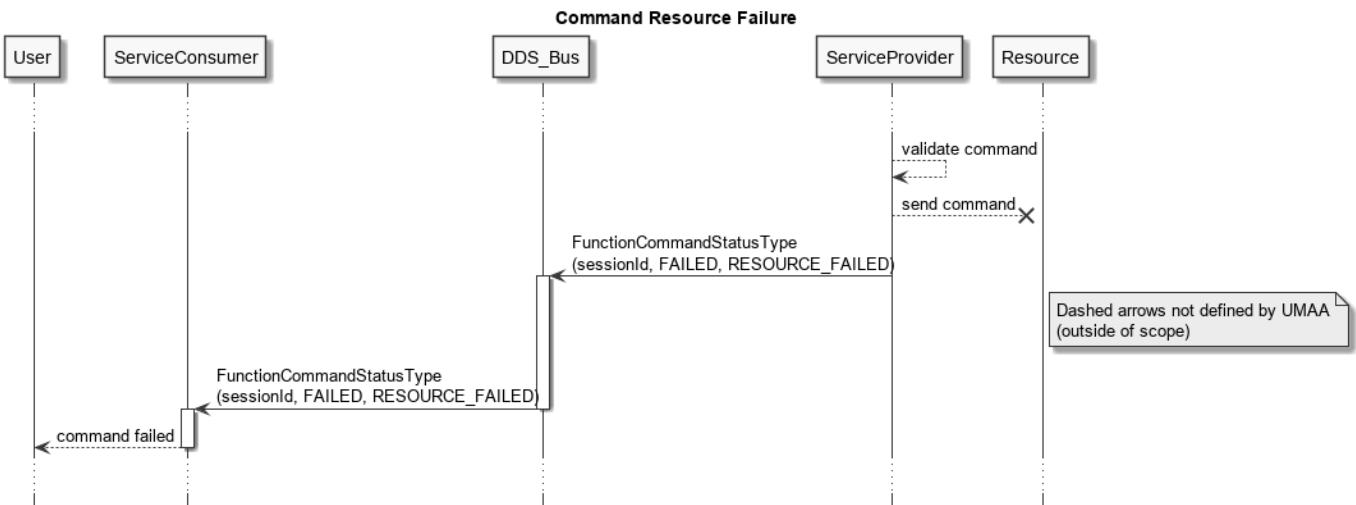


Figure 24: Sequence Diagram for a Command That Fails due to Resource Failure.

In the second example, the Resource takes too long to respond, so the Service Provider cancels the request and reports a `FunctionCommandStatusType` with a status of `FAILED` and a reason of `TIMEOUT`. This is shown in Figure 25.




Figure 25: Sequence Diagram for a Command That Times Out Before Completing.

Other failure conditions will follow a similar pattern: when the failure is recognized, the Service Provider will publish a `FunctionCommandStatusType` with a status of `FAILED` and a reason that reflect the cause of the failure.

4.1.4.5 Command Canceled The Service Consumer may decide to cancel the command before processing is finished. To signal a desire to cancel a command, the Service Consumer disposes of the existing `FunctionCommandType` from the DDS bus before the execution is complete. When notified of the command disposal, and if the Service Provider is able to cancel the command, it should respond to the Service Consumer with a `FunctionCommandStatusType` with both the status and reason as `CANCELED`. At this point, the DDS bus should dispose of the `FunctionCommandStatusType`, the `FunctionCommandAckReportType` and, (if defined for the Function service) the `FunctionExecutionStatusReportType`. This is shown in Figure 26. If the command cannot be canceled, then the Service Provider can continue to update the command status until the execution is completed. Reporting will include `FunctionCommandStatusType` with a status of `COMPLETED` and a reason of `SUCCEEDED`. Then, the DDS bus should dispose of the `FunctionCommandStatusType`, the `FunctionCommandAckReportType`, and (if defined for the Function service) the `FunctionExecutionStatusReportType`.

There is no new, unique, or specific status message response to a cancel command from the Service Provider. The cancel command status can be inferred through the corresponding `FunctionCommandStatusType` status and reason updates.

On loss of liveness of a Service Provider while executing a command, all Service Consumers must cancel (dispose) all in-process commands with that Service Provider.

On loss of liveness of a Service Consumer while executing a command, all Service Providers must treat the command as canceled. This means the service should report the `CANCELED` status for the command, and then dispose the command status, ack, and execution status (if one exists).




Figure 26: Sequence Diagram for a Command That is Canceled by the Service Consumer Before the Service Provider can Complete It.

4.1.5 Command Cleanup

The Service Consumer and Service Provider are responsible for disposing of corresponding data that is published to the DDS bus when the command is no longer active. With the exception of a canceled command, the signal that a **FunctionCommandType** can be disposed is when the **FunctionCommandStatusType** reports a terminal state (**COMPLETED** or **FAILED**)³. In turn, the signal that a **FunctionCommandStatusType**, **FunctionCommandAckReportType**, and (if defined for the Function service) the **FunctionExecutionStatusReportType** can be disposed is when the corresponding **FunctionCommandType** has been disposed. This is shown in Figure 27.

³While **CANCELED** is also a terminal state, the **CANCELED** command cleanup is handled specially as part of the cancelling sequence and, as such, does not need to be handled here.




Figure 27: Sequence Diagram Showing Cleanup of the Bus When a Command Has Been Completed and the Service Consumer No Longer Wishes to Maintain the Commanded State.

4.1.6 Command Shutdown Sequence

As part of shutdown, both the Service Provider and Service Consumer are required to perform a shutdown sequence. This shutdown cleans up resources on the DDS bus and informs the system that the Service Provider and Service Consumer are no longer available.

The Service Provider and Service Consumer can shut down in any order. The sequence diagram is shown in Figure 28.




Figure 28: Sequence Diagram for Command Shutdown.

4.1.6.1 Service Provider Shutdown Sequence During shutdown, the Service Provider is required to fail any incomplete requests and then unregisters as a publisher of the `FunctionCommandStatusType`, `FunctionCommandAckReportType`, and (if defined for the Function service) the `FunctionExecutionStatusReportType`.

The Service Provider is also required to unsubscribe from the `FunctionCommandType`.

The Service Provider Shutdown sequence is shown in Figure 29.




Figure 29: Sequence Diagram for Command Shutdown for Service Providers.

4.1.6.2 Service Consumer Shutdown Sequence During shutdown, the Service Consumer is required to cancel any incomplete requests and then unregister as a publisher of the `FunctionCommandType`.

The Service Consumer is also required to unsubscribe from the `FunctionCommandStatusType`, the `FunctionCommandAckReportType` if subscribed, and the `FunctionExecutionStatusReportType` if defined for the Function service and subscribed.

The Service Consumer Shutdown sequence is shown in Figure 30.




Figure 30: Sequence Diagram for Command Shutdown for Service Consumers.

4.2 Request / Reply

This section defines the flow of control for request/reply over the DDS bus. A request/reply is used to obtain data or status from a specific Service Provider.

A Service Provider is required to reply to all requests it receives. In the case of requests with no query data, this is accomplished via a DDS subscribe. In the case of a request with associated query data, a message with the query data must be published by the requester. To direct a request at a specific Service Provider or set of services, UMAA defines a destination GUID as part of requests.

The sequence diagrams in Sections 31 through 35 demonstrate different exchanges between a Service Consumer and Service Provider. Within the diagrams, the dashed arrows represent implementation-specific communications that are outside of UMAA's scope. Additionally, these sequence diagrams are examples of one possible implementation. Other implementations may have different communication patterns between the Service Provider and the Resource, or be implemented completely within the Service Provider process itself (no external Resource). However, in all implementations, UMAA-defined exchanges with the DDS bus between the Service Consumer and Service Provider must happen in the order shown within the sequence diagrams.

4.2.1 Request/Reply without Query Data

Figure 31 shows the sequence of exchanges in the case where there is no specific query data (i.e., the service is always just providing the current data to the bus).

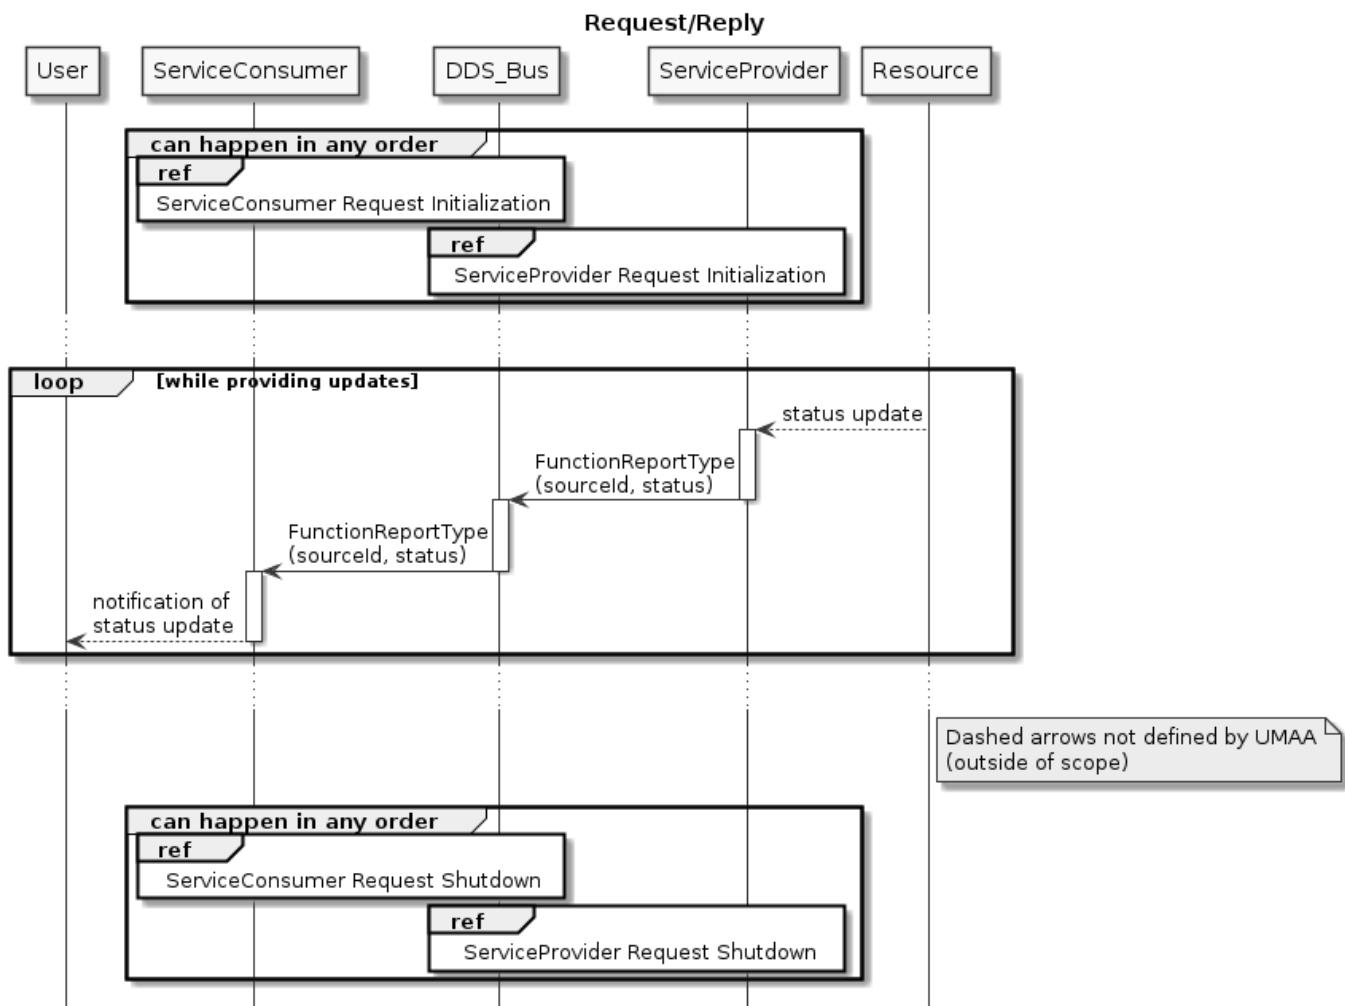


Figure 31: Sequence Diagram for a Request/Reply for Report Data That Does Not Require any Specific Query Data.

4.2.1.1 Service Provider Startup Sequence The Service Provider registers as a publisher of `FunctionReportTypes` to be able to respond to requests. The Service Provider must also handle reports that exist on the bus from a previous instantiation, either by providing an immediate update or, if the status is unrecoverable, disposing of the old `FunctionReportType`. This is shown in Figure 32.

As `FunctionReportType` updates are required (either through event-driven changes or periodic updates), the Service Provider publishes the updated data. The DDS bus will deliver the updates to the Service Consumer.




Figure 32: Sequence Diagram for Initialization of a Service Provider to Provide FunctionReportTypes.

4.2.1.2 Service Consumer Startup Sequence The Service Consumer subscribes to the FunctionReportType to signal an outstanding request for updates. This is shown in Figure 33.




Figure 33: Sequence Diagram for Initialization of a Service Consumer to Request FunctionReportTypes.

4.2.1.3 Service Provider Shutdown To no longer provide FunctionReportTypes, the Service Provider disposes of the FunctionReportType and unregisters as a publisher of the data (shown in Figure 34).




Figure 34: Sequence Diagram for Shutdown of a Service Provider.

4.2.1.4 Service Consumer Shutdown To no longer request FunctionReportTypes, the Service Consumer unsubscribes from FunctionReportType (shown in Figure 35).



Figure 35: Sequence Diagram for Shutdown of a Service Consumer.

4.2.2 Request/Reply with Query Data

Currently, UMAA does not define any request/reply interactions with query data, but it is expected that some will be defined. When defined, this section will be expanded to describe how they must be used.

5 Support Operations (SO) Services and Interfaces

5.1 Services and Interfaces

The interfaces in the following subsections describe how each UCS-UMAA topic is defined by listing the name, namespace, and member attributes. The "name" corresponds with the message name of a given service interface. The "namespace" defines the scope of the "name" where similar commands are grouped together. The "member attributes" are fields that can be populated with differing data types, e.g. a generic "depth" attribute could be populated with a double data value. Note that using a UCS-UMAA "Topic Name" requires using the fully-qualified namespace plus the topic name.

Each interface topic is referenced by a UMAA service and is defined as either an input or output interface.

Attributes ending in one or more asterisk(s) denote the following:

* = Key (annotated with @key in IDL file; vendors may use different notation to indicate a key field)

† = Optional (annotated with @optional in IDL file; vendors may use different notation to indicate an optional field)

Optional fields should be handled as described in the UMAA Compliance Specification.

Commands issued on the DDS bus must be treated as if they are immutable in UMAA and, therefore, if updated (treated incorrectly as mutable), the resulting service actions are indeterminate and flow control protocols are no longer guaranteed.

Operations without DDS Topics

⊕ = Operations that are handled directly in DDS

query<...> - All query operations are used to retrieve the correlated report message. For UMAA, this operation is accomplished through subscribing to the appropriate DDS topic.

cancel<...> - All cancel operations are used to nullify the current command. For UMAA, this operation is accomplished through the DDS dispose action on the publisher.

report<...>CancelCommandStatus - All cancel reports are included here to show completeness of the MDE model mapping to UMAA. For UMAA, this operation is not used. Instead, the cancel status is inferred from the associated command status. If the cancel command is successful, the corresponding command will fail with a command status and reason of CANCELED. If the corresponding command status reports COMPLETED, then this cancel command has failed.

5.1.1 HealthReport

The purpose of this service is to provide health details, which includes, but is not limited to, the health as determined from BITs (Built-In-Tests).

Table 8: HealthReport Operations

Service Requests (Inputs)	Service Responses (Outputs)
queryHealth⊕	reportHealth

See [Section 5.1](#) for an explanation of the inputs and outputs marked with a ⊕.

5.1.1.1 reportHealth

Description: This operation is used to report the most recent health status for each resource. This service is expected to report all managed resource/code pairs at all times (including reporting "NONE" severity when "No error condition exists"). Updates are required for each change event and at a configurable rate.

Namespace: UMAA::SO::HealthReport

Topic: HealthReportType

Data Type: HealthReportType

Table 9: HealthReportType Message Definition

Attribute Name	Attribute Type	Attribute Description
Additional fields included from UMAA::UMAAStatus		
logTime	DateTime	Log time when the error occurs.
severity	ErrorConditionEnumType	The type of error reported.
status†	StringLongDescription	A detailed, human-readable string which specifies the status of the system or subsystem, such as the reason for failure. Systems should not parse or use any information from this for processing purposes.
code*	ErrorCodeEnumType	The types of system or subsystems associated with the error report.
resourceID*	IdentifierType	Unique Identifier of the health detail of the resource.

5.1.2 LogReport

The purpose of this service is to provide log messages.

Table 10: LogReport Operations

Service Requests (Inputs)	Service Responses (Outputs)
queryLogReport⊕	reportLogReport

See [Section 5.1](#) for an explanation of the inputs and outputs marked with a ⊕.

5.1.2.1 reportLogReport

Description: This operation is used to report an entry on the bus that can be subscribed to and potentially entered into a system level log. This data is in addition to existing UMAA messages. It is a way of providing more detailed information than is normally published on the UMAA bus by the standard messaging.

Namespace: UMAA::SO::LogReport

Topic: LogReportType

Data Type: LogReportType

Table 11: LogReportType Message Definition

Attribute Name	Attribute Type	Attribute Description
Additional fields included from UMAA::UMAAStatus		

Attribute Name	Attribute Type	Attribute Description
entry	StringLongDescription	A human-readable description which specifies the contents of the log message. Systems should not parse or use any information from this for processing purposes.
level	LogLevelEnumType	Specifies the log level of the message.

5.2 Common Data Types

Common data types define DDS types that are referenced throughout the UMAA model. These DDS types are considered common because they can be re-used as the data type for many attributes defined in service interface topics, interface topics, and other common data types. These data types are not intended to be directly published to/subscribed as DDS topics.

5.2.1 UCSMDEInterfaceSet

Namespace: UMAA::UCSMDEInterfaceSet

Description: Defines the common UCSMDE Interface Set Message Fields.

Table 12: UCSMDEInterfaceSet Structure Definition

Attribute Name	Attribute Type	Attribute Description
timeStamp	DateTime	The origination time of the data being conveyed in the message, or as close to the data or command generation time as is reasonably possible.

5.2.2 UMAACommand

Namespace: UMAA::UMAACommand

Description: Defines the common UMAA Command Message Fields.

Table 13: UMAACommand Structure Definition

Attribute Name	Attribute Type	Attribute Description
Additional fields included from UMAA::UCSMDEInterfaceSet		
source*	IdentifierType	The unique identifier of the originating source of the command interface.
destination*	IdentifierType	The unique identifier of the destination of the command interface.
sessionID*	NumericGUID	The unique identifier for the session.

5.2.3 UMAAStatus

Namespace: UMAA::UMAAStatus

Description: Defines the common UMAA Status Message Fields.

Table 14: UMAAStatus Structure Definition

Attribute Name	Attribute Type	Attribute Description
Additional fields included from UMAA::UCSMDEInterfaceSet		
source*	IdentifierType	The unique identifier of the originating source of the status interface.

5.2.4 UMAACommandStatusBase

Namespace: UMAA::UMAACommandStatusBase

Description: Defines the common UMAA Command Status Base Message Fields.

Table 15: UMAACommandStatusBase Structure Definition

Attribute Name	Attribute Type	Attribute Description
Additional fields included from UMAA::UCSMDEInterfaceSet		
source*	IdentifierType	The unique identifier of the originating source of the command status interface.
sessionID*	NumericGUID	The unique identifier for the session.

5.2.5 UMAACommandStatus

Namespace: UMAA::UMAACommandStatus

Description: Defines the common UMAA Command Status Message Fields.

Table 16: UMAACommandStatus Structure Definition

Attribute Name	Attribute Type	Attribute Description
Additional fields included from UMAA::UMAACommandStatusBase		
commandStatus	CommandStatusEnumType	The status of the command.
commandStatusReason	CommandStatusReasonEnumType	The reason for the status of the command.
logMessage	StringLongDescription	Human-readable description related to response. Systems should not parse or use any information from this for processing purposes.

5.2.6 DateTime

Namespace: UMAA::Common::Measurement::DateTime

Description: Describes an absolute time. Conforms with POSIX time standard (IEEE Std 1003.1-2017) epoch reference point of January 1st, 1970 00:00:00 UTC.

Table 17: DateTime Structure Definition

Attribute Name	Attribute Type	Attribute Description
seconds	DateTimeSeconds	The number of seconds offset from the standard POSIX (IEEE Std 1003.1-2017) epoch reference point of January 1st, 1970 00:00:00 UTC.
nanoseconds	DateTimeNanoSeconds	The number of nanoseconds elapsed within the current DateTimeSecond.

5.2.7 IdentifierType

Namespace: UMAA::Common::IdentifierType

Description: This structure defines a two-level hierarchical identifier, where the parent is defined to be a group or collection of entities.

Table 18: IdentifierType Structure Definition

Attribute Name	Attribute Type	Attribute Description
id	NumericGUID	Provides the identifier of an entity.
parentID	NumericGUID	Provides the identifier of the parent, which is a group or collection of one or more entities. If the entity has no parent (it is the root of the tree), this value will be the Nil UUID.

5.3 Enumerations

Enumerations are used extensively throughout UMAA. This section lists the values associated with each enumeration defined in UCS-UMAA.

5.3.1 CommandStatusReasonEnumType

Namespace: UMAA::Common::MaritimeEnumeration::CommandStatusReasonEnumType

Description: Defines a mutually exclusive set of reasons why a command status state transition has occurred.

Table 19: CommandStatusReasonEnumType Enumeration

Enumeration Value	Description
CANCELED	Indicates a transition to the CANCELED state when the command is canceled successfully.
INTERRUPTED	Indicates a transition to the FAILED state when the command has been interrupted by a higher priority process.
OBJECTIVE_FAILED	Indicates a transition to the FAILED state when the commanded resource is unable to achieve the command's objective due to external factors.
RESOURCE_FAILED	Indicates a transition to the FAILED state when the commanded resource is unable to achieve the command's objective due to resource or platform failure.
RESOURCE_REJECTED	Indicates a transition to the FAILED state when the commanded resource rejects the command for some reason.
SERVICE_FAILED	Indicates a transition to the FAILED state when the commanded resource is unable to achieve the command's objective due to processing failure.
SUCCEEDED	Indicates the conditions to proceed to this state have been met and a normal state transition has occurred.
TIMEOUT	Indicates a transition to the FAILED state when the command is not acknowledged within some defined time bound.
UPDATED	Indicates a transition back to the ISSUED state from a non-terminal state when the command has been updated.
VALIDATION_FAILED	Indicates a transition to the FAILED state when the command contains missing, out-of-bounds, or otherwise invalid parameters.

5.3.2 ErrorCodeEnumType

Namespace: UMAA::Common::MaritimeEnumeration::ErrorCodeEnumType

Description: A mutually exclusive set of values that defines the error codes.

Table 20: ErrorCodeEnumType Enumeration

Enumeration Value	Description
ACTUATOR	Actuator
FILESYS	File system
NONE	None
POWER	Power
PROCESSOR	Processor
RAM	RAM

Enumeration Value	Description
ROM	ROM
SENSOR	Sensor
SOFTWARE	Software

5.3.3 ErrorConditionEnumType

Namespace: UMAA::Common::MaritimeEnumeration::ErrorConditionEnumType

Description: A mutually exclusive set of values that defines the error condition.

Table 21: ErrorConditionEnumType Enumeration

Enumeration Value	Description
ERROR	An error condition is reported and expected to seriously compromise use of the reporting component or device.
FAIL	An error condition is reported with severity indicating component or device failure.
INFO	An error condition is reported, but impact on operation and performance is minimal.
NONE	No error condition exists.
WARN	An error condition is reported and expected to have significant impact on component or device performance.

5.3.4 LogLevelEnumType

Namespace: UMAA::Common::MaritimeEnumeration::LogLevelEnumType

Description: Defines a mutually exclusive set of values that defines the log level.

Table 22: LogLevelEnumType Enumeration

Enumeration Value	Description
ERROR	An error message level.
INFORMATION	An informational message level.
WARNING	A warning message level.

5.3.5 CommandStatusEnumType

Namespace: UMAA::Common::MaritimeEnumeration::CommandStatusEnumType

Description: Defines a mutually exclusive set of values that defines the states of a command as it progresses towards completion.

Table 23: CommandStatusEnumType Enumeration

Enumeration Value	Description
CANCELED	The command was canceled by the requestor before the command completed successfully.
COMMANDED	The command has been placed in the resource's command queue but has not yet been accepted.
COMPLETED	The command has been completed successfully.
EXECUTING	The command is being performed by the resource and has not yet been completed.
FAILED	The command has been attempted, but was not successful.
ISSUED	The command has been issued to the resource (typically a sensor or streaming device), but processing has not yet commenced.

5.4 Type Definitions

This section describes the type definitions for UMAA. The table below lists how UMAA defined types are mapped to the DDS primitive types.

Table 24: Type Definitions

Type Name	Primitive Type	Range of Values	Description
DateTimeNanoseconds	long	units=Nanoseconds minInclusive=0 maxInclusive=999999999	The number of nanoseconds elapsed within the current second.
DateTimeSeconds	longlong	units=Seconds minInclusive=-9223372036854775807 maxInclusive=9223372036854775807	The seconds offset from the standard POSIX (IEEE Std 1003.1-2017) epoch reference point of January 1st, 1970 00:00:00 UTC.
LargeCollectionSize	long	maxInclusive=2147483647 minInclusive=0	Specifies the size of a Large Collection.
NumericGUID	octet[16]	minInclusive=0 maxInclusive=(2^{128})-1	Represents a 128-bit number according to RFC 4122 variant 2.
StringLongDescription	string	length=4095	Represents a long format description.

A Appendices

A.1 Glossary

Note: This glossary aims to define terms that are uncommon, or have a special meaning in the context of UMAA and/or the DoD. This glossary covers the complete UMAA specification. Not every word defined here appears in every ICD.

Almanac Data (GPS)	A navigation message that contains information about the time and status of the entire satellite constellation.
Coulomb	The SI unit of electric charge, equal to the quantity of electricity conveyed in one second by a current of one ampere.
Ephemeris Data (GPS)	A navigation message used to calculate the position of each satellite in orbit.
Glowplug or Glow Plug	A heating device used to aid in starting diesel engines.
Interoperability	1) The ability to act together coherently, effectively, and efficiently to achieve tactical, operational, and strategic objectives. 2) The condition achieved among communications-electronics systems or items of communications-electronics equipment when information or services can be exchanged directly and satisfactorily between them and/or their users.
Mean Sea Level	The average height of the surface of the sea for all stages of the tide; used as a reference for elevations.
Middleware	A type of computer software that provides services to software applications beyond those available from the operating system. Middleware makes it easier for software developers to implement communication and input/output, so they can focus on the specific purpose of their application.
SoaML	The Service oriented architecture Modeling Language (SoaML) specification that provides a metamodel and a UML profile for the specification and design of services within a service-oriented architecture. The specification is managed by the Object Management Group (OMG).

A.2 Acronyms

Note: This acronym list is included in every ICD and covers the complete UMAA specification. Not every acronym appears in every ICD.

ADD	Architecture Design Description
AGL	Above Sea Level
ASF	Above Sea Floor
BSL	Below Sea Level
BWL	Beam at Waterline
C2	Command and Control
CMD	Command
CO	Comms Operations
CPA	Closest Point of Approach
CTD	Conductivity, Temperature and Depth
DDS	Data Distribution Service
DTED	Digital Terrain Elevation Data
EGM	Earth Gravity Model
EO	Engineering Operations
FB	Feedback
GUID	Globally Unique Identifier
HM&E	Hull, Mechanical, & Electrical

ICD	Interface Control Document
ID	Identifier
IDL	Interface Definition Language Specification
IMO	International Maritime Organization
INU	Inertial Navigation Unit
LDM	Logical Data Model
LOA	Length Over All
LRC	Long Range Cruise
LWL	Length at Waterline
MDE	Maritime Domain Extensions
MEC	Maximum Endurance Cruise
MM	Mission Management
MMSI	Maritime Mobile Service Identity
MO	Maneuver Operations
MRC	Maximum Range Cruise
MSL	Mean Sea Level
OMG	Object Management Group
PIM	Platform Independent Model
PMC	Primary Mission Control
PNT	Precision Navigation and Timing
PO	Processing Operations
PSM	Platform Specific Model
RMS	Root-Mean-Square
ROC	Risk of Collision
RPM	Revolutions per minute
RTPS	Real Time Publish Subscribe
RTSP	Real Time Streaming Protocol
SA	Situational Awareness
SEM	Sensor and Effector Management
SO	Support Operations
SoaML	Service-oriented architecture Modeling Language
STP	Standard Temperature and Pressure
UCS	Unmanned Systems Control Segment
UMAA	Unmanned Maritime Autonomy Architecture
UML	Unified Modeling Language
UMS	Unmanned Maritime System
UMV	Unmanned Maritime Vehicle
UxS	Unmanned System
WGS84	Global Coordinate System
WMM	World Magnetic Model
WMO	World Meteorological Organization